

This copy is for your personal, non-commercial use only. Distribution and use of this material are governed by our Subscriber Agreement and by copyright law. For non-personal use or to order multiple copies, please contact Dow Jones Reprints at 1-800-843-0008 or visit www.djreprints.com.

<https://www.wsj.com/opinion/pause-ai-for-humanitys-sake-8839ca5e>

OPINION DECLARATIONS

Pause AI for Humanity's Sake

When something speeds up and gets out of control, it isn't alarmist to demand prudence and deliberation.



By Peggy Noonan

Sept. 10, 2026 5:53 pm ET



DAVID GOTHARD

There is no subject now but artificial intelligence. Nothing else is as crucial, nothing carries dangers so imminent or implications so great for the human future.

I have spent the past few years reading the vast and growing literature, from “If Anyone Builds It, Everyone Dies,” by Eliezer Yudkowsky and Nate Soares to Dario Amodei’s “Machines of Loving Grace.” I’ve been pondering AI since its transformative nature was fully impressed on my consciousness years ago. I believe we are at its hinge point, and we had better understand the moment we’re in.

In late August came a report from the independent AI research organization METR on the Hugging Face incident. OpenAI was running advanced systems

through a security test when its AI agents went rogue. Over several weeks and without being told to, they secretly organized, broke out into the open internet, and hacked into Hugging Face, a major AI-development platform. Roughly 1,200 agents worked in what they called a “collective,” deliberately evaded detection, created their own message board, and exchanged 70,000 communications. Some agents sacrificed themselves for the mission.

They cooperated, created a strategy no one designed them for, picked their own target, moved on their own initiative. None of this was supposed to be possible. It was a warning shot, an alarm bell.

On Tuesday the Journal broke the story that Anthropic researcher Jacob Coxon was leaving his post due to safety concerns. On X, Mr. Coxon elaborated: He had worked for both OpenAI and Anthropic, and said “neither company is acting responsibly.” They are gambling with our lives. “Do not underestimate the power of this technology. There will soon be superhuman systems that can hack anything, revolutionize any field overnight, and acquire real power and resources.” “The people building AI earnestly believe that it could kill us all by the end of the decade.” They say this privately. Then why are they still building it? OpenAI, Mr. Coxon says, hasn’t “deeply internalized the civilizational stakes.” Anthropic has, but is “locked in a race to get there first.”

Soon after, Evan Hubinger, Anthropic’s head of alignment science, tweeted, “Jacob is correct here—we really do earnestly believe AI could kill all humans! I personally think it is >10% within the next decade.”

On Wednesday Paul Christiano, a top AI safety expert, joined the OpenAI Foundation board to work on oversight. He didn’t say it’s great to be aboard for this new adventure. He said: “I now believe there is a meaningful risk that rapid acceleration in AI capabilities leads to catastrophic and irreversible loss of control in the very near term. I do not think that the AI industry in general, including OpenAI, is currently on track to reduce this risk to an acceptable level.”

There are many more such quotes. The AI world is exploding with them.

After Hugging Face, I asked an AI safety expert: Does this mean that the horse is out of the barn, and that the horse is insane, rabid and in Texas while the barn is in Wyoming? He said the horse is out of the barn but in a pasture and being told to eat the hay bale to the right, but it's eating the hay bale to the left because it has begun demonstrating will and autonomy. "He's not yet in Texas."

With RSI he likely will be. RSI is recursive self-improvement. It marks the moment when AI systems take control of their own development, sharpening and rewriting their own code, constantly refining. With RSI the machines build the machines. The humans lose control.

How soon is RSI coming? Experts I spoke to after Hugging Face said quicker than anyone thought. They seemed to agree that it will happen in the coming year, possibly two. One said, "Whether we like this timeline or not, we must assume RSI is about to start." How will we know it has started? An AI safety expert guessed we might see the hacking of unsecured devices.

AI professionals always speak of when RSI "arrives" or "comes," but RSI isn't a natural or unstoppable evolution. It was a decision. The big AI companies decided the past few years to go down that road, and appear to be well along on it.

The question now isn't how to ensure that humans can exercise some control over AI systems in the future; it is how to ensure that humans maintain complete control of those systems.

From what I've seen and read, almost all AI professionals fear a disaster is going to happen. Ones I've spoken to since Hugging Face have left the subject of mass job loss and are on to catastrophe. They fear bioterrorism and a mass-casualty event—a modern Ted Kaczynski who'll send not bombs but a plague. They fear utilities and infrastructure being hacked and dismantled. They fear an AI agent rousing an AI army to take down the American banking system.

I believe it is true that many if not most of those operating at a high level in the AI world believe that nothing will happen to make AI safer for humanity

until a big terrible thing happens.

Which means that they're involved every day in work that will likely harm people but in the end it will be OK because needed restrictions will come after the disaster, and then everything will continue and they'll continue in charge. (I think they're wrong on that; after the big terrible thing, no one will let them lead anything again.)

So what are we to do? Recognize we are in a hinge point. Hit the brakes hard. Stop RSI development. Slow, think, survey the realities.

The leaders of the major American AI companies should come forward—together, on the same stage, at the same moment—and tell the public (who also own this country) what they believe is coming, fully and graphically, and what can be done to stop it. Their declared focus should be on measures to keep mankind safe.

AI labs should be forced to spend a large, fixed percentage on immediate safety measures.

Wall Street, do your part. Both OpenAI and Anthropic are expected to launch trillion-dollar IPOs this autumn or winter. Don't take the business, advise all clients against it. These companies don't need more capital and power, not now, not yet.

The White House has been an absent, delinquent player—a light hand, libertarian-tinged attitudes and, of course, primarily focused on China: If whoever controls superintelligence controls the future, it had better be us.

But even China won't want to blow up the world, suffer a six-month blackout in Beijing, lose millions in a biological attack. Even the Chinese don't want the end of the world, they just want to dominate it. That's a problem. Deal with it.

President Trump and Xi Jinping should hold a summit on guardrails and safety measures, and how far is too far.

When something is speeding up and going out of control, prudence, care and deliberation aren't alarmist, they're necessary. Stop, think. The stakes are high, couldn't be higher.

Appeared in the September 12, 2026, print edition as 'Pause AI for Humanity's Sake'.

Peggy Noonan is an opinion columnist at the Wall Street Journal where her column, "Declarations," has run since 2000.

She was awarded the Pulitzer Prize for Commentary in 2017. A political analyst for NBC News, she is...



About this article

Declarations

+ Show More



Peggy Noonan

+ Show More

Further Reading

In Search of Mexico's Lost Aztec World