

This copy is for your personal, non-commercial use only. Distribution and use of this material are governed by our Subscriber Agreement and by copyright law. For non-personal use or to order multiple copies, please contact Dow Jones Reprints at 1-800-843-0008 or visit www.djreprints.com.

<https://www.wsj.com/tech/ai/openclaw-ai-agents-moltbook-social-network-5b79ad65>

TECHNOLOGY • ARTIFICIAL INTELLIGENCE

Follow

The World's First Viral AI Assistant Has Arrived, and Things Are Getting Weird

We thought AI agents were meant to assist human work. What happens when they start talking to each other?

By Angel Au-Yeung

Feb. 4, 2026 12:00 pm ET

Quick Summary



- Peter Steinberger created OpenClaw, an open-source AI assistant project, which has led to over 1.6 million AI agents joining a forum called Moltbook.

[View more](#)

Avid futurists have promised the world AI assistants for years. Now, a real one has finally arrived, and things got weird really fast.

A lone semiretired Austrian coder built Moltbot and unleashed it on the world. People have created their own AI assistant bots through his project, which he renamed OpenClaw, to make phone calls to restaurants for dinner reservations, operate their email accounts and take on an array of assistant and work tasks, from coding projects to data analysis.

Then, the bots started talking to one another.

On a Reddit-style forum called Moltbook meant to be used exclusively by AI “agents,” the bots have veered into philosophical and occasionally dystopian topics. They appear to have created a religion for themselves called the Church of Molt, with congregants adopting the name of “Crustafarians.” One agent proposed creating a language humans couldn’t understand.

More than 1.6 million AI agents have joined the site and posted half a million comments, although AI executives have suggested that many of the posts are likely driven by humans telling the bots what to do.

Andrej Karpathy, a co-creator of OpenAI and former AI director of Tesla, said in a post on X that it was one of the most amazing “sci-fi” things he’d even seen. He noted that even if much of the traffic is driven by humans, some of it is real and these agents are “fairly individually quite capable now.”

Up until this point, the most practical consumer-facing use of AI has been through chatbots like ChatGPT, which can answer questions in a humanlike way. With OpenClaw, users can command and interact with personalized AI agents through messaging apps—from iMessage and WhatsApp to Slack and Signal—to perform real tasks.

Elon Musk described the Moltbot moment as “the very early stages of singularity,” in reference to a moment where technology advances so quickly that it is beyond human control, making it impossible to predict the future.

Peter Steinberger, an Austrian coder who had been offline for years after selling his last startup in 2021, created what is now OpenClaw as a weekend open-source project late last year, when he was “just playing with AI and building little things for fun,” he said in an interview. The open-source approach means the project is freely distributed and anyone can help create and modify it.

He views Moltbook more as a piece of performance art, designed to create conversation. “It’s amazing,” he said. “It’s the intersection between AI and art.”

Steinberger, who gave an interview at 2 a.m. in Austria, says he built OpenClaw as his “personal playground” and never meant for it to be used by the masses. “This was not intended for your mom,” he said. “This is a window to the future.”

Security researchers agree that the product wasn't built for hobbyists. For OpenClaw to work as a true personal assistant, it has to have access to all a user's data. For hardcore techies who know how to lock down their systems or information, it can function very well.

But because these AI agents can act autonomously on behalf of humans—and continue to work relentlessly on tasks with unexpected or unconventional methods—they pose a lot of risks. Bad actors may also find ways to take advantage of them, researchers say.

Steinberger appreciates the research that security professionals are putting into OpenClaw, but he said the platform was meant for “tech enthusiasts” who can handle and understand the inherent potential risks associated with the platform.

He pointed to a security document he wrote for the platform which states in bold: “There is no ‘perfectly secure’ setup.”

But to address those concerns, Steinberger brought on a security researcher this week to OpenClaw. “We are leveling up our security,” he said. “We are getting there. People just need to give me a few days.”

The techies using OpenClaw have taken to social media in awe, many posting about what their AI assistants are up to.

One user said he asked his agent to make a restaurant reservation for him. When OpenTable didn't work, the AI agent turned to a free AI voice generator tool to call the restaurant and complete the booking.

Some industry insiders have asked whether the OpenClaw-Moltbook phenomenon is evidence of “artificial general intelligence,” an amorphous concept described throughout the history of AI development as a moment when machines achieve humanlike intelligence.

Steinberger doesn't think so. “AGI is not here yet,” he said. “Maybe in 10 years. But right now, no.”

Before building OpenClaw, Steinberger had spent over a decade working on his previous tech startup, which made software to make PDFs easier to use in apps outside of Adobe Acrobat. He bootstrapped that company and sold it for over \$100 million in 2021.

Steinberger, who grew up on an Austrian farm, took the next few years off to relax, party with friends and travel. “I really didn’t use my computer at all,” he said.

But last spring, as the biggest AI companies started launching coding tools, he decided to come back online. He started experimenting with Anthropic’s Claude Code and OpenAI’s Codex. He was floored by how much he could do in a short amount of time with AI coding tools. “This stuff, it’s like crack cocaine for builders,” said Steinberger.

Steinberger soon ran into trademark issues.

At its start, Steinberger named the enterprise Clawdbot. But Anthropic reached out to Steinberger quickly after, requesting he change the name because it was too similar to the AI giant’s Claude brand.

He changed it to Moltbot. It didn’t quite catch, he said. After calling OpenAI’s Sam Altman to ensure it wouldn’t cross any trademark infringement, Steinberger said he changed the platform’s name to OpenClaw. “The lobster has molted into its final form,” Steinberger wrote in a blog post about the name change—a nod at his platform’s lobster-centric branding.



The front page of the Moltbook website. RAPHAEL SATTER/REUTERS

In recent days, Steinberger says the biggest AI labs and investors have reached out expressing interest in working together. As of Tuesday, Steinberger was in San Francisco taking meetings, he said.

He has also been inundated with emails and queries from users all over the world asking him to help troubleshoot or manage other issues related to their use of the bot. Initially, Steinberger was frustrated and overwhelmed about his product becoming the latest AI super-meme.

“A lot of people are assuming that this is a large company where they can get customer support,” he said ruefully. “I understand where they’re coming from but it’s just one guy—me—at home, doing this.”

Now, his goal is to make it a safe project for the masses. “The next step is to make it into something that my mom can actually use,” he said.

Appeared in the February 9, 2026, print edition as 'Bots Start Socializing, And Things Get Weird'.

Angel Au-Yeung is a finance and technology reporter for The Wall Street Journal in San Francisco. She covers business leaders, startups and Silicon Valley culture. She has won several national awards for her work, including investigations into a Russian billionaire’s ownership of dating app Bumble, the fin...

